

Cloud Computing final project

A cloud-based file storage system
and the OSU benchmarks in k8s

Isac Pasianotto

2024-03-21

Cloud **basic** module

- 1 Cloud **basic** module
 - Requirements
 - Nextcloud: the chosen solution
 - Deployed structure
 - Load test: Locust
 - Scalability
- 2 Cloud **advanced** module:
Nextcloud in a kubernetes flavor
- 3 Cloud **advanced** module:
OSU benchmarks in k8s

Requirements

Identify, deploy and implement a cloud-based file storage system

- User / admin roles
- User should be able to login, upload, download, delete files, have a personal folder
- Admin should be able to manage users, quotas, etc
- Security evaluation
- Scalability evaluation
- Test the system with a load test





- Comes with already built-in all the required features (User/admin, download/upload, private folder, etc)
- Has a large community and a lot of plugins
- Can be easily deployed using docker (already available image)
- Supports different types of storage (local, NFS, S3, etc)
- Supports various database backends (MySQL, PostgreSQL, SQLite)

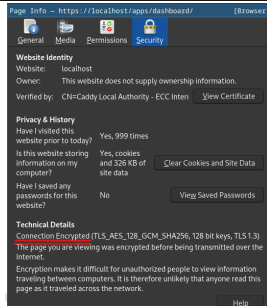
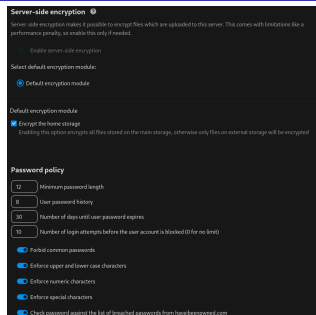
Nextcloud: the chosen solution (con'd)

- Possibility to enable **server-side encryption**^a
- Customizable **password minimum requirements and password policy**
- **Brute-force protection** (additional app)
- **Two-factor authentication**
- Accepts **OAuth2** external authentication

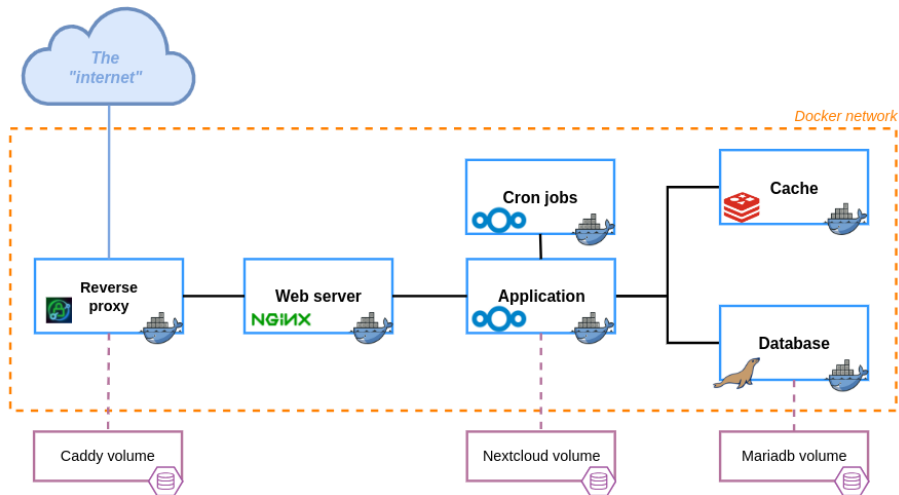
And what about the connection user-server?

⇒ **Caddy** automatically provides a **Let's Encrypt** certificate

^aDrawback: more CPU usage



Deployed structure



Load test: Locust


Locust is an open-source load testing tool, which allows defining user behavior with Python code.

- start with 1 user
- add 1 user every second until a predefined quota is reached
- every second, each active user tries to perform a random task
- each user has a wait time between 0.5 and 3 seconds before declaring the task as failed



Tasks	%
upload_small	30%
upload_medium	35%
upload_large	5%
download	30%

Load test: Locust (con'd)



Load test performed with server-side encryption enabled. During the test, I got warnings from the server about the CPU usage over 90%.

I've considered all the possible solutions in the case of a real-world scenario

- **On-site** solution: Build on a owned server, with appropriate hardware
- **Infrastructure as a Service** (IaaS): rent needed resources from a cloud provider, but manage everything by yourself
- **Platform as a Service** (PaaS): rent a pre-configured environment, and focus only on the application and data
- **Software as a Service** (SaaS): rent a ready-to-use application, just use it with your data

Scalability (con'd)

	<i>on-site</i>	<i>IaaS</i>	<i>PaaS</i>	<i>SaaS</i>
hardware acquisition/ cost	high	null	null	null
maintenance costs	high	low	null	null
predictable costs	high	medium	low	low
Renting costs	null	low	medium	high
Need to care about security	high	medium	on your app	trust the vendor
Physically own the data	yes	no	no	no
Scalability	completely on you	rent more resources	pay more	pay more
Time to get ready	high	medium	low	almost null

Cloud **advanced** module: Nextcloud in a kubernetes flavor

Cloud **advanced** module:

Nextcloud in a kubernetes flavor

- 1 Cloud **basic** module
- 2 Cloud **advanced** module:
Nextcloud in a kubernetes flavor
 - Requirements
 - Different architecture
 - How it works
 - Benefits of kubernetes
- 3 Cloud **advanced** module:
OSU benchmarks in k8s

Requirements

Re-deploy the solution of the previous exercise, but in a kubernetes environment

- The cluster must run `k8s`, one node is required
- The volume must survive a `pod crash` and `accidental deletion`
- The service must be accessible from the user via IP or FQDN
- Eventual databases or **third-party services** necessary for the deployed software must **run in their pod**

Different architecture

- One node → the control plane is also the worker node
- Used the nextcloud [helm chart](#)
- Used [metallb](#) to expose the service
- Access is done through a [service](#)
- All the pods run in a dedicated namespace
- The kubernetes orchestrator automatically takes care of the pods, respawning them if needed

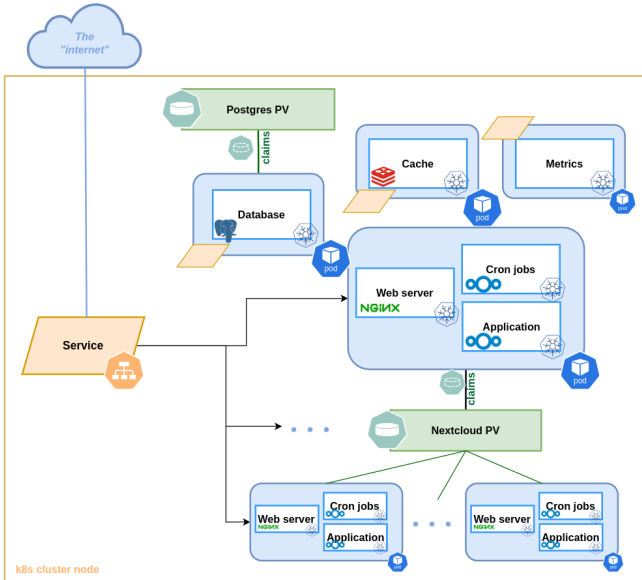


mariadb.existingSecret doesn't seem to be working #506



v1nsai opened this issue on Jan 4 · 7 comments

Different architecture (con'd)



How it works

- Installation done with **helm** and a custom **values.yaml** file
- Manually created **persistent volume** and **persistent volume claim**
- Manually created **secrets** for storing credentials (avoiding them to be stored in the **values.yaml** file)
- Postgres, Redis and metrics run in their own pods
- System can survive a pod crash or accidental deletion

```
Context: kubernetes-admin/kubernetes
Cluster: kubernetes
User: kubernetes-admin
K8s Rev: v0.20.2
K8s Rev: v0.20.7
CPU: 4/4
MEM: 4/4
```

NAME	PF	READY	RESTARTS	STATUS	IP	NODE	AGE
my-nextcloud-instance-677d8648-6sx9b	●	1/3	0	Terminating	10.17.6.15	es2-00	14s
my-nextcloud-instance-677d8648-rgh2t	●	2/3	0	Running	10.17.6.18	es2-00	28s
my-nextcloud-instance-677d8648-sjwzr	●	3/3	0	Running	10.17.6.14	es2-00	14s
my-nextcloud-instance-677d8648-vp2bt	●	2/3	0	Running	10.17.6.12	es2-00	14s
my-nextcloud-instance-metrics-859c6996d-nbsk7	●	1/1	0	Running	10.17.6.12	es2-00	14s
my-nextcloud-instance-postgresql-0	●	1/1	0	Running	10.17.6.17	es2-00	14s
my-nextcloud-instance-redis-master-0	●	1/1	0	Running	10.17.6.16	es2-00	14s

After a pod deletion, the orchestrator automatically respawns it

```
vagrant@es2-00 nextcloud-helm$ kubectl get service -n nextcloud
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nextcloud-instance	LoadBalancer	10.102.241.179	192.168.121.286	8080:32744/TCP	12m
my-nextcloud-instance-metrics	LoadBalancer	10.96.133.84	192.168.121.281	9293:31843/TCP	12m
my-nextcloud-instance-postgresql	ClusterIP	10.97.249.153	<none>	5432/TCP	12m
my-nextcloud-instance-postgresql-hl	ClusterIP	None	<none>	5432/TCP	12m
my-nextcloud-instance-redis-headless	ClusterIP	None	<none>	6379/TCP	12m
my-nextcloud-instance-redis-master	ClusterIP	10.101.68.84	<none>	6379/TCP	12m

```
vagrant@es2-00 nextcloud-helm$ curl -s 192.168.121.286:8080 | head -n 15
```

```
<!DOCTYPE html>
<html class="ng-csp" data-placeholder-focus="false" lang="en" data-locale="en" translate="no" >
  <head>
    <data-requesttoken="vd9b0WKSiosT2uf2TM3DxkuRQq4FntvMQLU25u6Hpb8=5LS/agvVP3t37aPL6P61lHHEupx3Zln9HKLDrgvVMxy=">
    <meta charset="utf-8">
    <title>
      Nextcloud
    </title>
    <meta name="viewport" content="width=device-width, initial-scale=1.0, minimum-scale=1.0">
    <meta name="apple-touch-icon" href="/core/img/favicons/apple-touch.png">
    <meta name="theme-color" content="#0082C9">
    <link rel="icon" href="/core/img/favicons/icon">
    <link rel="apple-touch-icon" href="/core/img/favicons/apple-touch.png">
    <link rel="mask-icon" sizes="any" href="/core/img/favicons/mask.svg" color="#0082C9">
    <link rel="manifest" href="/core/img/manifest.json" crossorigin="use-credentials">
    <link rel="stylesheet" href="/apps/theming/css/default.css?v=70e2b24f-0">
```

Benefits of kubernetes

Some of the benefits of using kubernetes as container orchestrator are:

- **Availability:** automatic pod respawning in case of crash
- **Scalability:** easy to add more nodes, and the [Horizontal Pod Autoscaler](#) can automatically add more pods
- **Topology awareness:** easy to schedule pods on different nodes to avoid single point of failure
- **Resource management:** kubernetes can automatically manage the resources of the pods, setting limits to the CPU and memory usage
- **Portability:** moving the configuration to another cluster/cloud provider should be easy

Cloud **advanced** module: OSU benchmarks in k8s

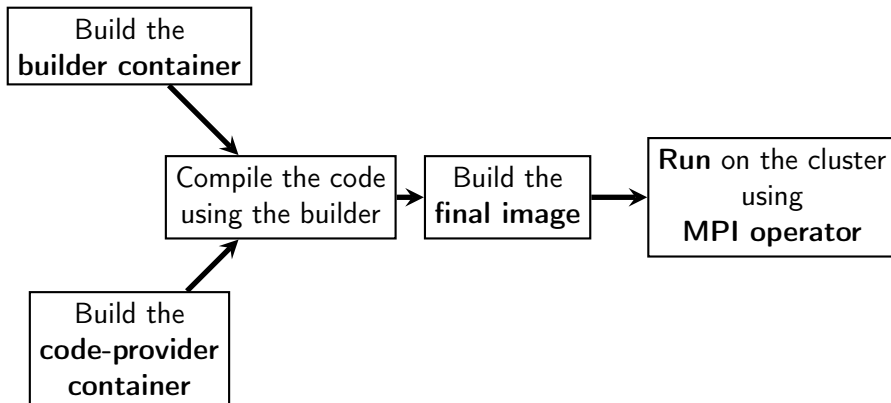
Cloud **advanced** module: OSU benchmarks in k8s

- 1 Cloud **basic** module
- 2 Cloud **advanced** module:
Nextcloud in a kubernetes flavor
- 3 Cloud **advanced** module:
OSU benchmarks in k8s
 - Requirements
 - Building the image
 - Calico vs Flannel
 - results

Requirements

- 2 nodes [k8s](#) cluster
- Nodes must talk with either [Calico](#) or [Flannel](#)
- [MPI](#) operator
- Perform the [OSU](#) benchmark to evaluate the latency between the nodes

Building the image


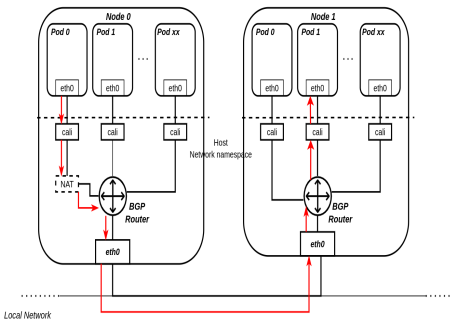


#174 · 00

No need to reinvent the wheel:

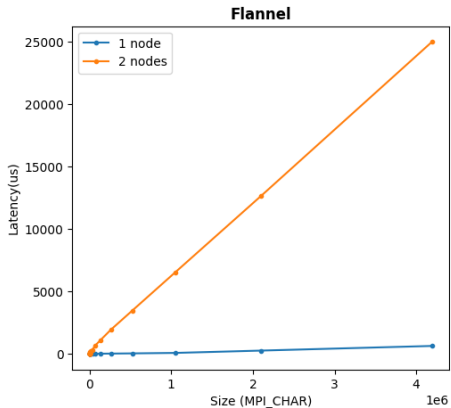
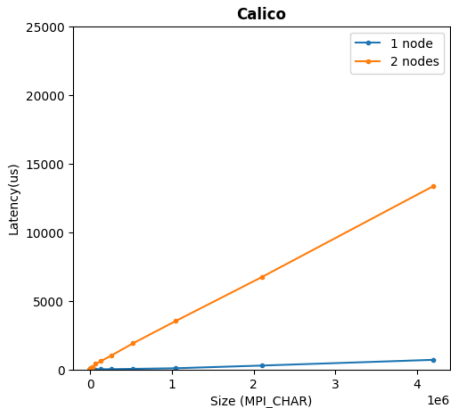
Calico vs Flannel

From pod 0 (node 0) to pod 1 (node 1):



Results¹


Latency 1 node vs 2 nodes




¹Each point represents the average of 20 runs.

Results (con'd)

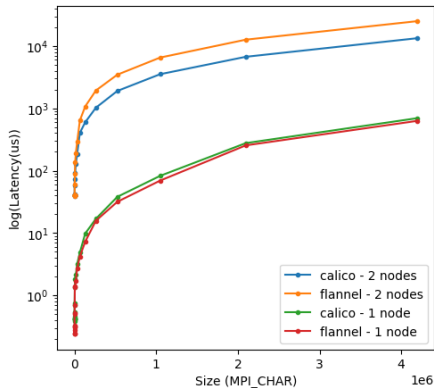
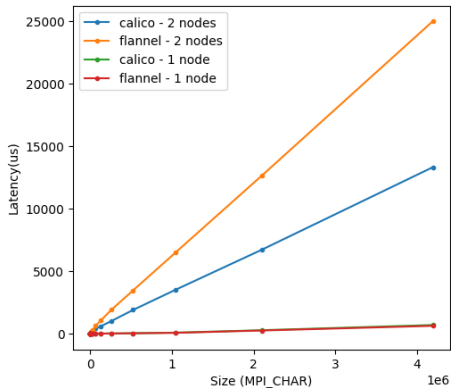
1 node



2 nodes



Results (con'd)



*Thank you for your
attention*